

Education

IIT, Kharagpur

🎓 M.Tech - Computer Sci.& Data Processing
📅 May 2019 📍 Kharagpur, India

Nagpur University

🎓 B.E. - Electronics & Power
📅 June 2015 📍 Chandrapur, India

Maharashtra State Board

🎓 12th Certificate - Science
📅 May 2011 📍 Chandrapur, India

Maharashtra State Board

🎓 10th Certificate
📅 May 2009 📍 Chandrapur, India

Skills

Languages & Frameworks

- Python, Pytorch, Huggingface
- Deepspeed, Accelerate
- FastAPI, MCP, A2A, LangGraph
- TRL, Axolotl, Llama Factory, Unsloth
- Docker, Kubernetes, Mlflow
- TGI, LoRAx, VLLM, Claude Code, Github Copilate

Cloud Services

- AWS (EC2, Sagemaker, S3, ECR, Inferntia2, Bedrock)
- GCP (Vertex AI, GSU)
- CloudNativx, Fireworks

Related Experience

Senior CV Developer - MLAI Technologies (4 Months)

Worked on POC for AI-driven KYC update automation to streamlined processes, cut costs, saved time, and boosted efficiency

Research Intern - TRDDC (4 Months)

Worked with TRDDC's NLP research team to develop an Answer Evaluation System.

Miscellaneous Projects

- Implemented scalable deployment of LLMs using Kubernetes orchestration combined with VLLM, enhancing model serving efficiency and enabling robust, distributed inference at scale.
- Trained and deployed a Cascade Mask R-CNN model using OpenMMLab's MMDetection framework for table extraction and document digitization, achieving an IoU score of 0.95.
- Develop a framework using Hugging Face and Ray to support real-time adapter switching and request batching
- Android UI Testing Agent using Multimodel LLM and appium server.

Experience

Senior Data Scientist || SIMPLIFYHEALTHCARE

📅 Oct 2025 – Present 📍 Pune, MH, IN (Hybrid)

- Working on agentic AI solutions in healthcare, building an agent orchestration framework and intelligent chatbots to improve user experience and operational efficiency.

Senior Software AI Engineer - 1 || GUPSHUP TECHNOLOGIES

📅 Dec 2021 – sept 2025 📍 Mumbai, MH, IN (Remote)

- Led the end-to-end training and inference of open-source LLMs on client / domain-specific data to improve task alignment. Identifying and implementing cost-effective solutions that maintain high throughput and low latency.
- Working on multi-agent orchestration frameworks to enable complex task automation and inter-agent collaboration.
- Collaborate cross-functionally with product, back-end, and infrastructure teams to integrate LLM-based systems into production pipelines.
- Manage and mentor a team of 2-3 full-time engineers and guide several interns, ensuring alignment with project milestones and innovation goals.

Associate Tech Lead - AI || ZYCUS

📅 Aug 2020 – Nov 2021 📍 Bengaluru, KA, IN (Remote)

- Enhanced Merlin, a legal document understanding framework, by integrating advanced AI capabilities, improving automation, accuracy, and overall document intelligence.

Machine Learning Engineer || QUANTIPHI ANALYTICS

📅 July 2019 – April 2020 📍 Mumbai, MH, IN (On-site)

- Developed and delivered ML/DL-based POCs to solve business problems.
- Evaluated models and frameworks to align with business needs and ensure effective solution delivery.

Recent Projects

Gupshup Agentic AI [In Development] || GUPSHUP AI TEAM

- Working on migration of existing microservices and FastAPI endpoints to MCP servers to support scalable agent integration.
- Working A2A protocol leveraging MCP servers to enable seamless agent communication and execution of intelligent tasks.
- Designing Multi-agent orchestration logic to support parallel and sequential task delegation across agents.

Tech Stack: MCP, A2A, FastAPI, Microservices, OpenAI, Gemini, Opensource Models

RAG Pipeline for ABB || GUPSHUP AI TEAM

- Designed and implemented a hybrid indexing strategy combining embedding and keyword-based retrieval techniques to optimize information access from diverse document sources.
- Integrated knowledge graph retrieval with Graphiti for structured reasoning and contextual understanding over entity linked information.
- Used RAGA for automated evaluation of retriever and generator performance.

Tech Stack: RAGA, Polars, FAISS, FastAPI, TFIDF, OpenAI, Gemini

Gupshup LLMs || GUPSHUP AI TEAM

- Build end-to-end ML pipeline for training, evaluation, and deployment.
- Utilized AWS S3 for efficient data storage and preprocessing, while training environments were containerized using AWS ECR and orchestrated with SageMaker for scalable compute.
- Leveraged DeepSpeed and Hugging Face Accelerate for efficient, cost-effective large-scale training. Deployed models using VLLM on a Kubernetes cluster for low-latency inference.
- Integrated MLflow for experiment tracking and model management, improving reproducibility and deployment speed.

Tech Stack: AWS, Deepspeed/Accelerate, LoRA, DPO, VLLM, Axolotl, Llama Factory